

Image Reduction & Classification

David Anderton and Ray Dedhia

Abstract—As the use of artificial intelligence continues to expand into the realm of image compression¹², questions are being raised on how well the original artefacts and nature of the source image are preserved. The consequences of being able to utilize smaller output images that maintain the majority of pertinent information from their often significantly larger source images are huge. With smaller inputs, applications from self driving cars that rely on camera input, autonomous flight, through to training neural nets could require significantly less processing time and resources. In this paper, we briefly examine how the Google Cloud Vision API responds to the content of a selection of images pre- and post-compression. We also consider a variety of compression techniques and extents. The goal: produce a broad estimate of how much an image can be compressed until the Google Cloud Vision API, as a proxy for other classification algorithms, can no longer determine useful insights about the constitution of the image.

An interactive exploration of the results is available at the following URI: <https://mit18065-dimensions.herokuapp.com/>

I. INTRODUCTION

In this paper, we briefly examine how the Google Cloud Vision API responds to the content of a selection of images pre- and post-compression. We also consider a variety of compression techniques and extents. The goal: produce a broad estimate of how much an image can be compressed until the Google Cloud Vision API, as a proxy for other classification algorithms, can no longer determine useful insights about the constitution of the image.

II. METHODOLOGY

Our methodology was as follows:

Within our study we utilised 22 unique images, from four major categories: six animals, six famous landmarks, five landscapes, five street signs. Expanding this to both the color and grayscale versions of these images, we had a resulting set of 44 source images.

After collating the source images we sent them through the Google Cloud Vision API³ and get a control set of labels for the 44 images. Then compress the images, using the five following methods: (1) Posterization through k-means clustering, (2) low-rank approximation using Principle Component Analysis (PCA), (3) “Mean pooling” operation commonly used in neural networks, (4) Mogrify (a command-line utility) and (5) Dropout (randomly removing a specific percentage of the pixels).

¹<https://ai.googleblog.com/2016/09/image-compression-with-neural-networks.html>

²<http://www.wave.one/icml2017/>

³The Google Cloud Vision API, <https://cloud.google.com/vision/>, classifies images into categories and detects objects and faces in images utilising non-public machine learning models.

Once the images were compressed in all of the above methods we then re-queried the Google Cloud Vision API to re-classify the images post-compression and considered the results.

III. COMPRESSION METHODS

A. Posterization

K-means clustering is an algorithm that groups a set of data points into k groups by mapping each data point to one of k centroids such that the distance between each data point and its centroid is minimized. It does this by randomly generating k centroids, (1) assigning each data point to the nearest centroid, (2) updating the centroid values by setting them equal to the arithmetic mean of the data points assigned to them in step 1, and repeating steps 1 and 2 until some stopping criteria has been met. The resulting k-means clustering implemented a posterization effect.

We arbitrarily stopped our k-means clustering algorithm after 10 iterations as an intuitive stopping point for efficiency, and used $k = 4$ and $k = 8$. In figure 1 below are the results of putting an image of a cat through 4-means and 8-means clustering.



Fig. 1: Image of a cat compressed using 4-means and 8-means clustering.

B. PCA

The principle components of some matrix, M , are its left and right singular vectors, u_i and v_i respectively, and its singular values, σ_i . The decomposition of a matrix into these components is referred to most commonly as the Singular Value Decomposition, or SVD. PCA is a form of matrix analysis that utilizes the largest principle components of a matrix of data in order to create a summarization or approximate the underlying data.

By utilising the Eckart-Young theorem⁴, the rank k matrix M_k closest to M is the sum of the product of the k largest singular values of M with their corresponding left and right singular vectors. That is, $M_k = \sum_{i=1}^k \sigma_i u_i v_i^T$.

⁴See Golub et al. <https://www.ime.usp.br/~jstern/miscellanea/seminario/Golub87.pdf>

Thus, PCA can be utilised to find the best low-rank approximation for M . Within our study we used PCA to reduce our image matrices to rank-5, rank-10 and rank-25 matrices. Figure 2 below shows the results of calculating the best rank-5, rank-10, and rank-25 approximations of an image of the Tokyo Tower.

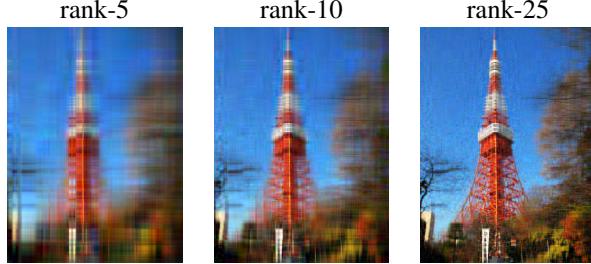


Fig. 2: Best rank-5, rank-10, and rank-25 matrix approximations of an image of the Tokyo Tower.

C. Mean Pooling

Mean pooling refers to the operation of dividing an image into regions of size $m \times n$, calculating the arithmetic mean of those regions, and outputting an image consisting of just those calculated averages.⁵

For our study we considered regions of size 16×16 , 32×32 and 64×64 . In figure 3 below are the results of putting an image of a US traffic light sign through a 16×16 , 32×32 , and 64×64 mean pooling filter.



Fig. 3: Image of a US traffic sign put through 16×16 , 32×32 , and 64×64 mean pooling filters.

D. Mogrify

The mogrify command line utility, part of the Image Magick suite of tools⁶, can be used to perform a number of operations on images, such as blurring, cropping, and compression. We used mogrify to compress our images using the option `-quality [num]`⁷. For JPEG image compression the expected value of `[num]` is an integer from 1 to 100. We used the value of 1, that applies the most extreme JPEG compression from mogrify. In figure 4 below, an image of the ocean before and after it has been

⁵See the UFLDL Tutorial on Pooling: <http://ufldl.stanford.edu/tutorial/supervised/Pooling/>

⁶See ImageMagick's documentation of the mogrify command: <https://www.imagemagick.org/script/mogrify.php>

⁷See ImageMagick's documentation of the quality option: <https://www.imagemagick.org/script/command-line-options.php#quality>

compressed by the command `mogrify -quality 1` can be considered.



Fig. 4: Image of the ocean that has been compressed with the command `mogrify -quality 1`.

E. Dropout

Inspired by the concept of dropout layers in neural networks⁸, which randomly remove nodes in the neural network to prevent overfitting, we implemented an algorithm that randomly removes percentage p of the pixels in an image. It does this by removing some percentage p of the columns and the same percentage p of the pixels from each of the remaining columns.

We implemented this algorithm with $p = 0.2, 0.5$, where p is a value between 0 and 1. In figure 5 the results of putting an image of a tiger through dropout with $p = 0.2$ and $p = 0.5$ can be considered.



Fig. 5: Image of a tiger put through dropout with $p = 0.2, 0.4$.

IV. CLASSIFICATION DATA

A. Before Compression

Google Cloud Vision assigned each image a set of labels. All of the labels assigned to the color images were accurate or close to accurate, and for each color image, at least one label closely identified the image (e.g. “tiger” for the iamge of a tiger and “tower” for the image of the Tokyo Tower).

However, not all the grayscale images were correctly classified. Ignoring extremely general labels such as “black and white” and “monochrome photography,” 40.9% of the grayscale images weren’t closely identified, and 18.2% of them were only given incorrect labels.

⁸See Srivastava et al.: <http://jmlr.org/papers/volume15/srivastava14a.old/srivastava14a.pdf>

B. After Compression

Color Images

method of compression	% images given at least one correct label	% images given best label
mogrify	81.8%	68.2%
dropout (20%)	77.3%	68.2%
dropout (50%)	59.1%	36.4%
pooling (16x16)	90.9%	63.6%
pooling (32x32)	40.9%	18.2%
pooling (64x64)	4.5%	0%
PCA (25)	95.5%	81.8%
PCA (10)	77.3%	40.9%
PCA (5)	27.3%	0%
clustering (8)	100%	90.9%
clustering (4)	86.4%	68.2%

Table 1: Analyzes how successfully color images were labeled by Google Cloud Vision after image compression.

Grayscale Images

method of compression	% images given at least one correct label	% images given best label
mogrify	59.1%	40.9%
dropout (20%)	17/22	15/22
dropout (50%)	13/22	8/22
pooling (16x16)	20/22	14/22
pooling (32x32)	9/22	4/22
pooling (64x64)	1/22	0/22
PCA (25)	21/22	18/22
PCA (10)	17/22	9/22
PCA (5)	6/22	0/22
clustering (8)	22/22	20/22
clustering (4)	19/22	15/22

Table 2: Analyzes how successfully grayscale images were labeled by Google Cloud Vision after image compression.

As expected, the less the image was compressed, the more accurate the labels it was assigned were. In addition, each form of image compression had its own set of issues.

The smoothening effect of posterization using the k-means clustering algorithm sometimes removed textures that Google Cloud Vision needed to correctly classify parts of an image. For instance, the image of the waterfall, after it was put through the 4-means and 8-means clustering algorithm, was assigned the label “snow” by Google Cloud Vision because Google Cloud Vision mistook the smoothed water for snow (see Fig. 6 below).



Fig. 6: Over-smoothening of an image of a waterfall as a result of the 8-means clustering algorithm.

Low-rank image matrix approximation using PCA often resulted in image distortion in a plaid-like pattern, especially in smaller images (see Fig. 7 below). This resulted in misclassifications such as Google Cloud Vision giving images in all four categories false labels such as “cloth,” “pattern,” and “line.”

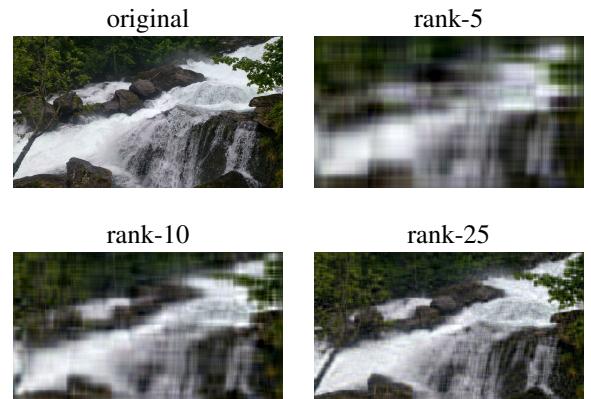


Fig. 7: Distortions in the low-rank matrix approximations of an image of a waterfall.

The “mean pooling” operation with a 64×64 filter, and, in some cases, the 32×32 filter as well, blurred smaller to the point that they become unrecognizable and Google Cloud Vision could no longer give them accurate labels (see Fig. 3). When the filter being used with the “mean pooling” operation is too large proportional to the size of the image, it removes all useful information from the image.

Image compression using the `mogrify -quality 1` command sometimes resulted in random bright colors showing up in the compressed images that did not occur in the

original images (see Fig. 8 below).



Fig. 8: Purple highlights appear in an image of mountains after it is compressed using the command mogrify -quality 1

The dropout algorithm we created resulted in distortion of the compressed images, almost as though the images were being reflected on moving water. This could potentially be fixed by adding a smoothening factor to our dropout algorithm, for example by using an algorithm such as seam-carving⁹

V. ANALYSIS

From the data, it appears that k-means clustering was the best compression method in terms of maintaining the legibility of the images for Google Cloud Vision. In a sense, this makes sense, because posterization generally simplifies images while retaining the overall shapes in the images.

In future analyses, it may be helpful to establish a more specific definition of accuracy and a way to compare the accuracy of Google Cloud Vision when classifying images before and after compression, as the table above does not consider a number of factors, such as the certainty of different labels, and the number of completely incorrect labels assigned to images.

VI. CONCLUSIONS

This information is useful because data reduction has been used as a method to fight adversarial attacks, in which stickers, random noise, etc. are added to images, causing neural networks to misclassify them.¹⁰ ¹¹ For example, putting adversarial stickers on stop signs can cause neural networks to be unable to detect stop signs.¹² The amount by which one can compress an image before it becomes illegible to neural networks, as well as the best method to use to compress an image to maintain legibility, can inform methods that use data reduction to fight adversarial attacks.

ACKNOWLEDGMENT

We would like to thank Prof. Strang for his support and interesting and informative lectures.

⁹For a description of seam-carving, see: <https://cs.brown.edu/courses/cs129/results/proj3/taox/>

¹⁰<https://openreview.net/pdf?id=S10qYwywf>

¹¹https://www.researchgate.net/publication/315890594_Dimensionality_Reduction_as_a_Defense_against_Evasion_Attacks_on_Machine_Learning_Classifiers

¹²<http://bair.berkeley.edu/blog/2017/12/30/yolo-attack/>

REFERENCES

- [1] Johnston, Nick and David Minnen. "Image Compression With Neural Networks." Google AI Blog. Published 29 September 2016. <https://ai.googleblog.com/2016/09/image-compression-with-neural-networks.html>
- [2] "Real-Time Adaptive Image Compression." WaveOne, Inc. WaveOne. Published May 2017. <http://www.wave.one/icml2017/>
- [3] "Cloud Vision API." Google Cloud. <https://cloud.google.com/vision/>
- [4] Golub, G. H., Alan Hoffman and G. W. Stewart. "A Generalization of the Eckart-Young-Mirsky Matrix Approximation Theorem." <https://www.ime.usp.br/~jstern/miscellanea/seminario/Golub87.pdf>
- [5] "Pooling." UFLDL Tutorial. <http://ufldl.stanford.edu/tutorial/supervised/Pooling/>
- [6] "mogrify." ImageMagick. <https://www.imagemagick.org/script/mogrify.php>
- [7] "command line options." ImageMagick. <https://www.imagemagick.org/script/command-line-options.php\#quality>
- [8] Srivastava, Nitish, et. al. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting." *Journal of Machine Learning Research*, 15 (2014): 1929-1958. Published 14 June 2014. <http://jmlr.org/papers/volume15/srivastava14a.old/srivastava14a.pdf>
- [9] Tao, Xiaofeng. "Seam Carving." <https://cs.brown.edu/courses/cs129/results/proj3/taox/>
- [10] Gopalakrishnan, Soorya, et. al. "Combating Adversarial Attacks Using Sparse Representations." ICLR 2018. <https://openreview.net/pdf?id=S10qYwywf>
- [11] Bhagoji, Arjun, et. al. "Dimensionality Reduction as a Defense against Evasion Attacks on Machine Learning Classifiers." https://www.researchgate.net/publication/315890594_Dimensionality_Reduction_as_a_Defense_against_Evasion_Attacks_on_Machine_Learning_Classifiers
- [12] Evtimov, Ivan, et. al. "Physical Adversarial Examples Against Deep Learning Networks." Berkeley Artificial Intelligence Research. Published 30 December 2017. <http://bair.berkeley.edu/blog/2017/12/30/yolo-attack/>